


Topic 2: Processes



Silberschatz, Galvin and Gagne ©2015


Rudowsky ©2005

Walpole ©2010

Kubiatowicz ©2010

Stallings ©2015

- Chapter 3: Process Concept
- Chapter 4: Multithreaded Programming
- Chapter 5: Process Scheduling



Contents

■ Processes


- Process concept
- Process states
- Process Control Block (PCB)
- Operations on Processes

■ Threads

- Thread concept
- Libraries to create threads

■ Operating System Examples


■ CPU scheduling





Objectives

- To introduce the notion of a process -- a program in execution, which forms the basis of all computation
- To describe the various features of processes, including scheduling, creation and termination, and communication
- To introduce the notion of a thread — a fundamental unit of CPU utilization that forms the basis of multithreaded computer systems
- To introduce CPU scheduling, which is the basis for multiprogrammed operating systems
- To describe various CPU-scheduling algorithms





Contents

■ Process


- Process concept
- Process states
- Process Control Block (PCB)
- Operations on Processes

■ Threads

- Thread concept
- Libraries to create threads

■ Operating System Examples


■ CPU scheduling





Process Concept


- An operating system executes a variety of programs
- **Program:**
 - description of how to perform an activity (algorithm)
 - instructions and static data values
 - static file (image)
- **Process:**
 - a program in execution; process execution must progress in sequential fashion
 - a snapshot of a program in execution
 - an instance of a program running in a computer





Process Concept


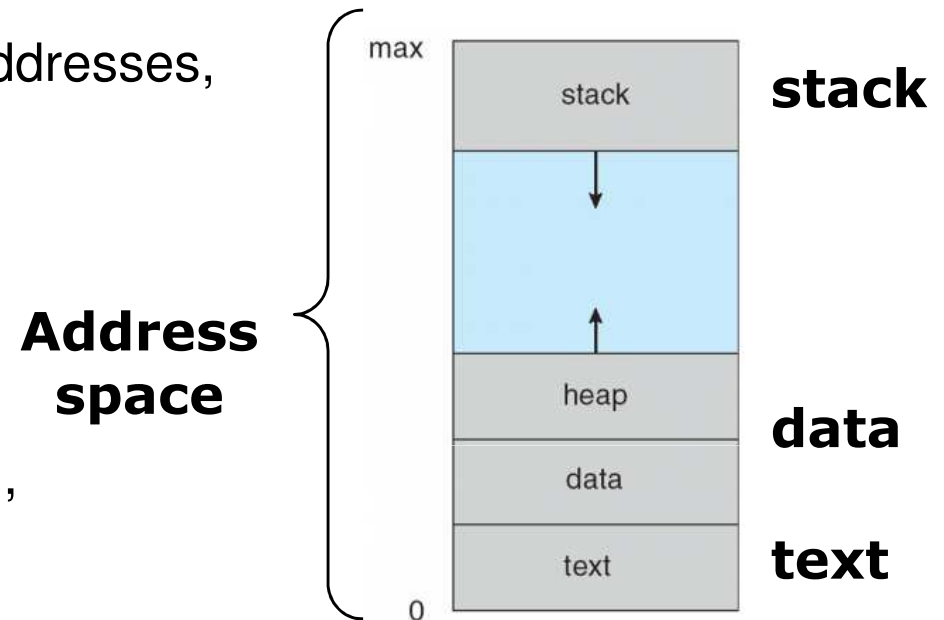
- A process is the basic unit of execution in an operating system
 - Each process has a number, its process identifier (pid)
- Different processes may run different instances of the same program





Process requirements

- At a minimum, process execution requires following resources:
 - Memory to contain the program code (**text section**) and data
 - **Stack** containing temporary data
 - ▶ Function parameters, return addresses, local variables
 - A set of CPU registers to support execution
 - ▶ CPU state (registers, Program Counter (PC), Stack Pointer (SP), etc)
 - ▶ operating system state (open files, accounting statistics etc)





Contents

■ Process


- Process concept
- Process states
- Process Control Block (PCB)
- Operations on Processes

■ Threads

- Thread concept
- Libraries to create threads

■ Operating System Examples


■ CPU scheduling






Process states

- As a process executes, it changes **state**
- Possible process states
 - **running**: instructions are being executed
 - **blocked (waiting)**: the process is waiting for some event to occur
 - **ready**: the process is waiting to be assigned to a processor




1. Process blocks for input
2. Scheduler picks another process
3. Scheduler picks this process
4. Input becomes available






Two-State Process Model

- Process may be in one of two states
 - ▶ Running
 - ▶ Not-running



(a) State transition diagram

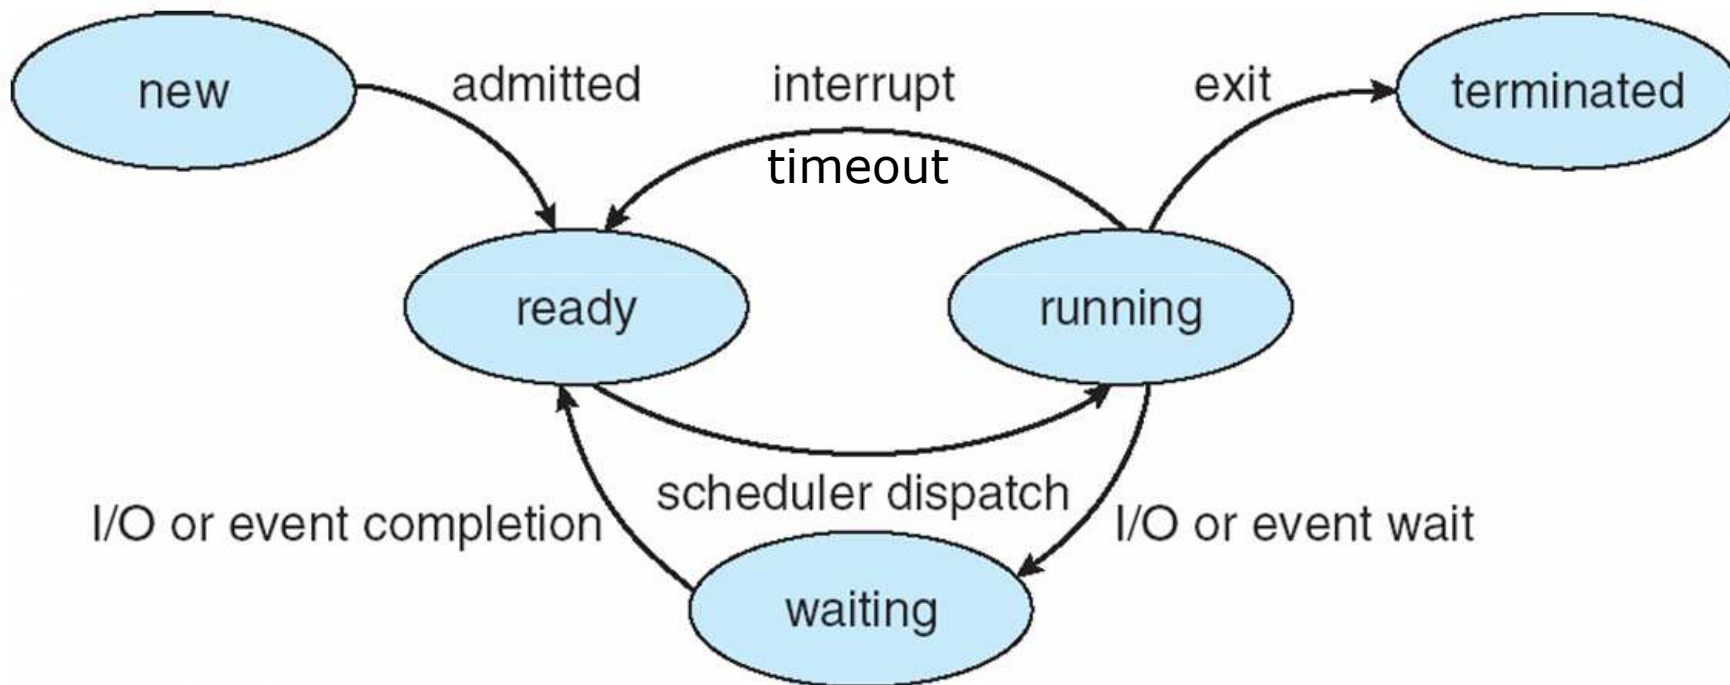




Five-State Process Model

■ States:


- **new**: The process is being created
- **running**: Instructions are being executed
- **waiting (blocked)**: The process is waiting for some event to occur
- **ready**: The process is waiting to be assigned to a processor
- **terminated**: The process has finished execution





Swapped and Suspended Processes


- Processor is faster than I/O so all processes could be waiting for I/O
 - Swap these processes to disk to free up more memory
 - ▶ More processes can be executed
- System overloaded
 - blocked state becomes suspend state when swapped to disk
- Two new states
 - Blocked/Suspend
 - Ready/Suspend





Reasons for Process Suspension

- Swapping: OS needs to release memory to bring in a ready process
- Protection: OS may suspend a process suspected of causing a problem
- User request: may wish to suspend execution for debugging reasons
- Timing: process executed periodically, and suspended while waiting the next round
- Others





Suspended Processes

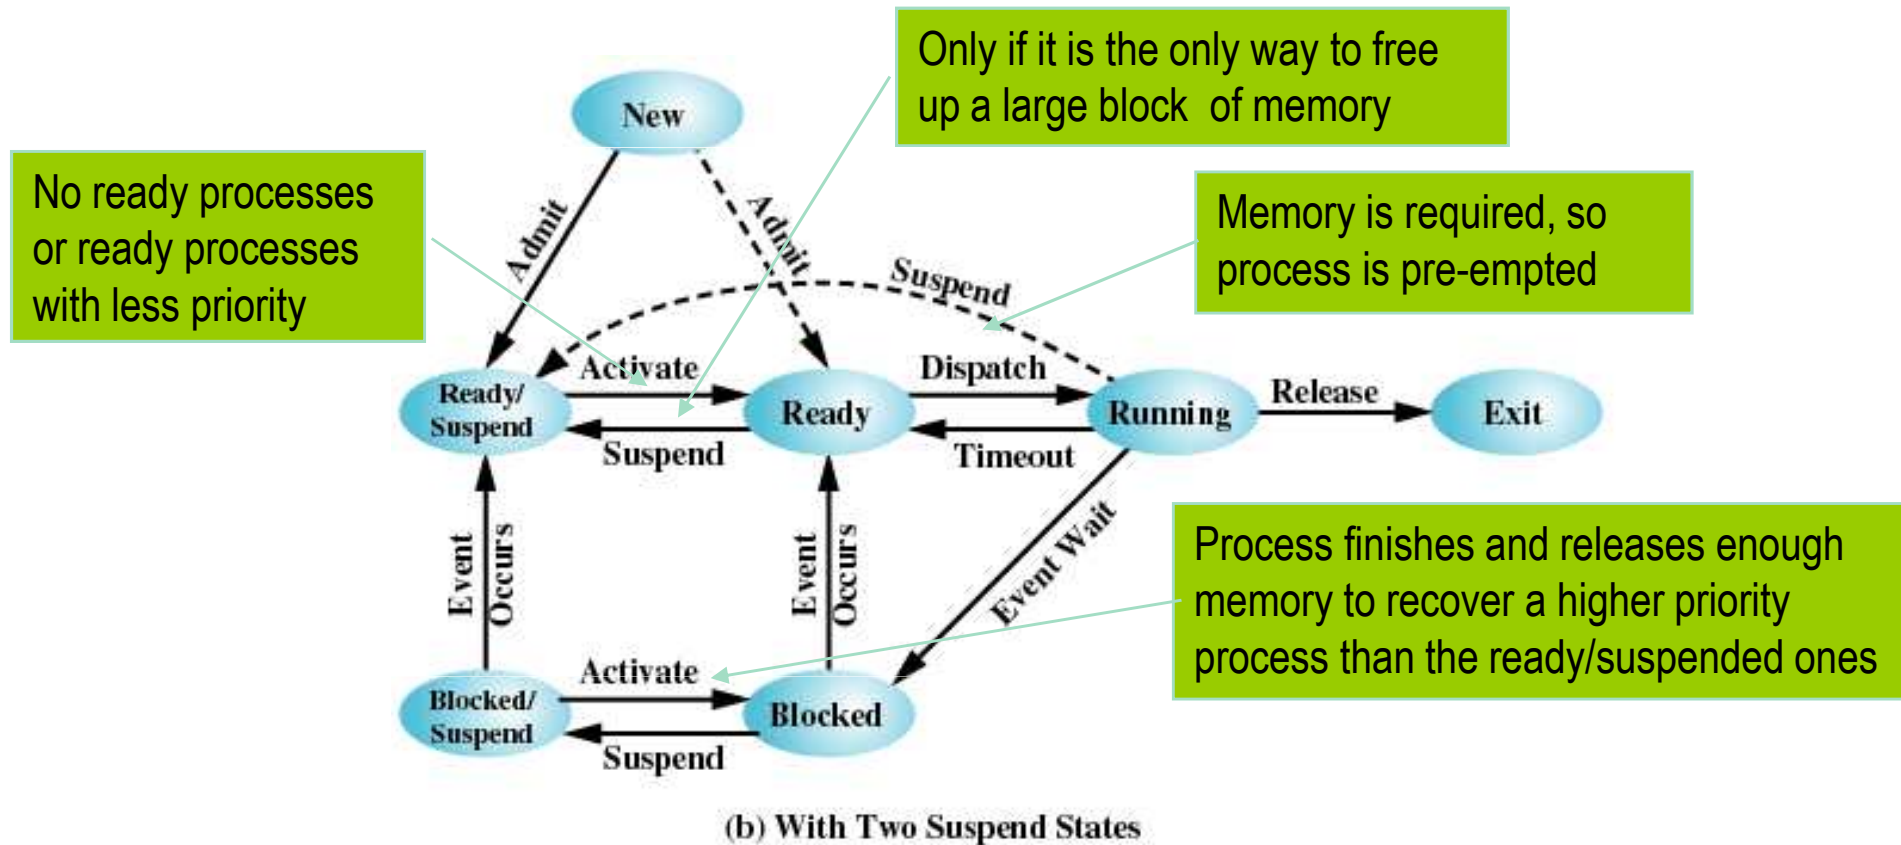




Figure 3.9 Process State Transition Diagram with Suspend States





Process Tracking



Source: Sistemas Operativos. Una visión aplicada.

a: A in CPU, B y C blocked
b: A calls the OS for I/O
c: All processes blocked (CPU idle)
d: B I/O finishes (awake and dispatch)
e: B in execution
f: C I/O finishes (awake), B ready
g: B still in CPU and C ready
h: B performs syscall. OS dispatches C and blocks B

i: C in CPU, A y B blocked
j: C calls the OS for I/O and A awakes
k: A in execution
l: I/O interrup. calls OS to wake up B
m: A still in Run and B watis for ready
n: A blocks
o: B starts to run

CPU always busy except in c



Contents

■ Process


- Process concept
- Process states
- Process Control Block (PCB)
- Operations on Processes


■ Threads

- Thread concept
- Libraries to create threads

■ Operating System Examples


■ CPU scheduling





Process Control Block

- What happens if the processor switches from process A to process B?
 - The processor needs information to get back to execute process A
 - Operating systems need a information structure






Process Control Block

- What happens if the processor switches from process A to process B?
 - The processor needs information to get back to execute process A
 - Operating systems need a information structure



Process Control Block (PCB)





Process Control Block


- What happens if the processor switches from process A to process B?
 - The processor needs information to get back to execute process A
 - Operating systems need a information structure



Process Control Block (PCB)



Context Switch





Process Control Block (PCB)

Information associated with **each process**
(also called **task control block**)

- **Process state** – running, waiting, etc
- **Program counter** – location of instruction to next execute
- **CPU registers** – contents of all process-centric registers
- **CPU scheduling information** – priorities, scheduling queue pointers
- **Memory**-management information – memory allocated to the process
- **Accounting information** – CPU used, clock time elapsed since start, time limits
- **I/O status information** – I/O devices allocated to process, list of open files

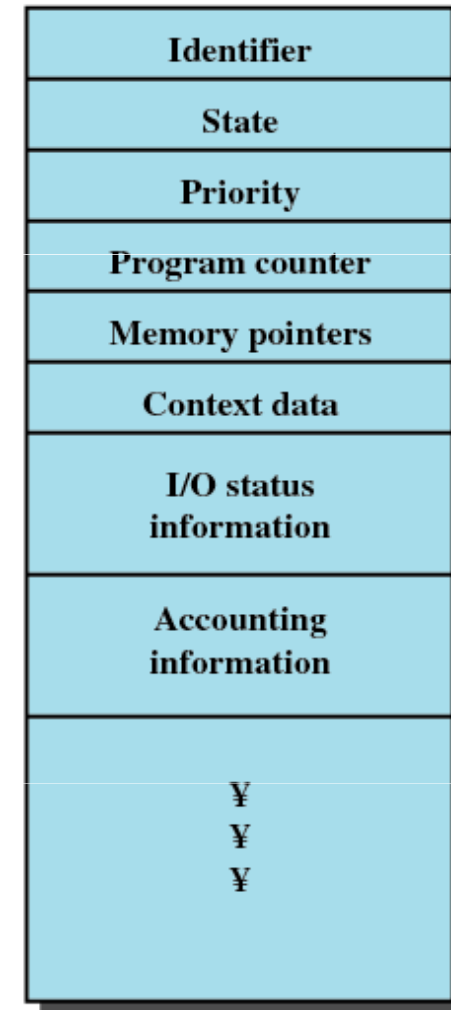
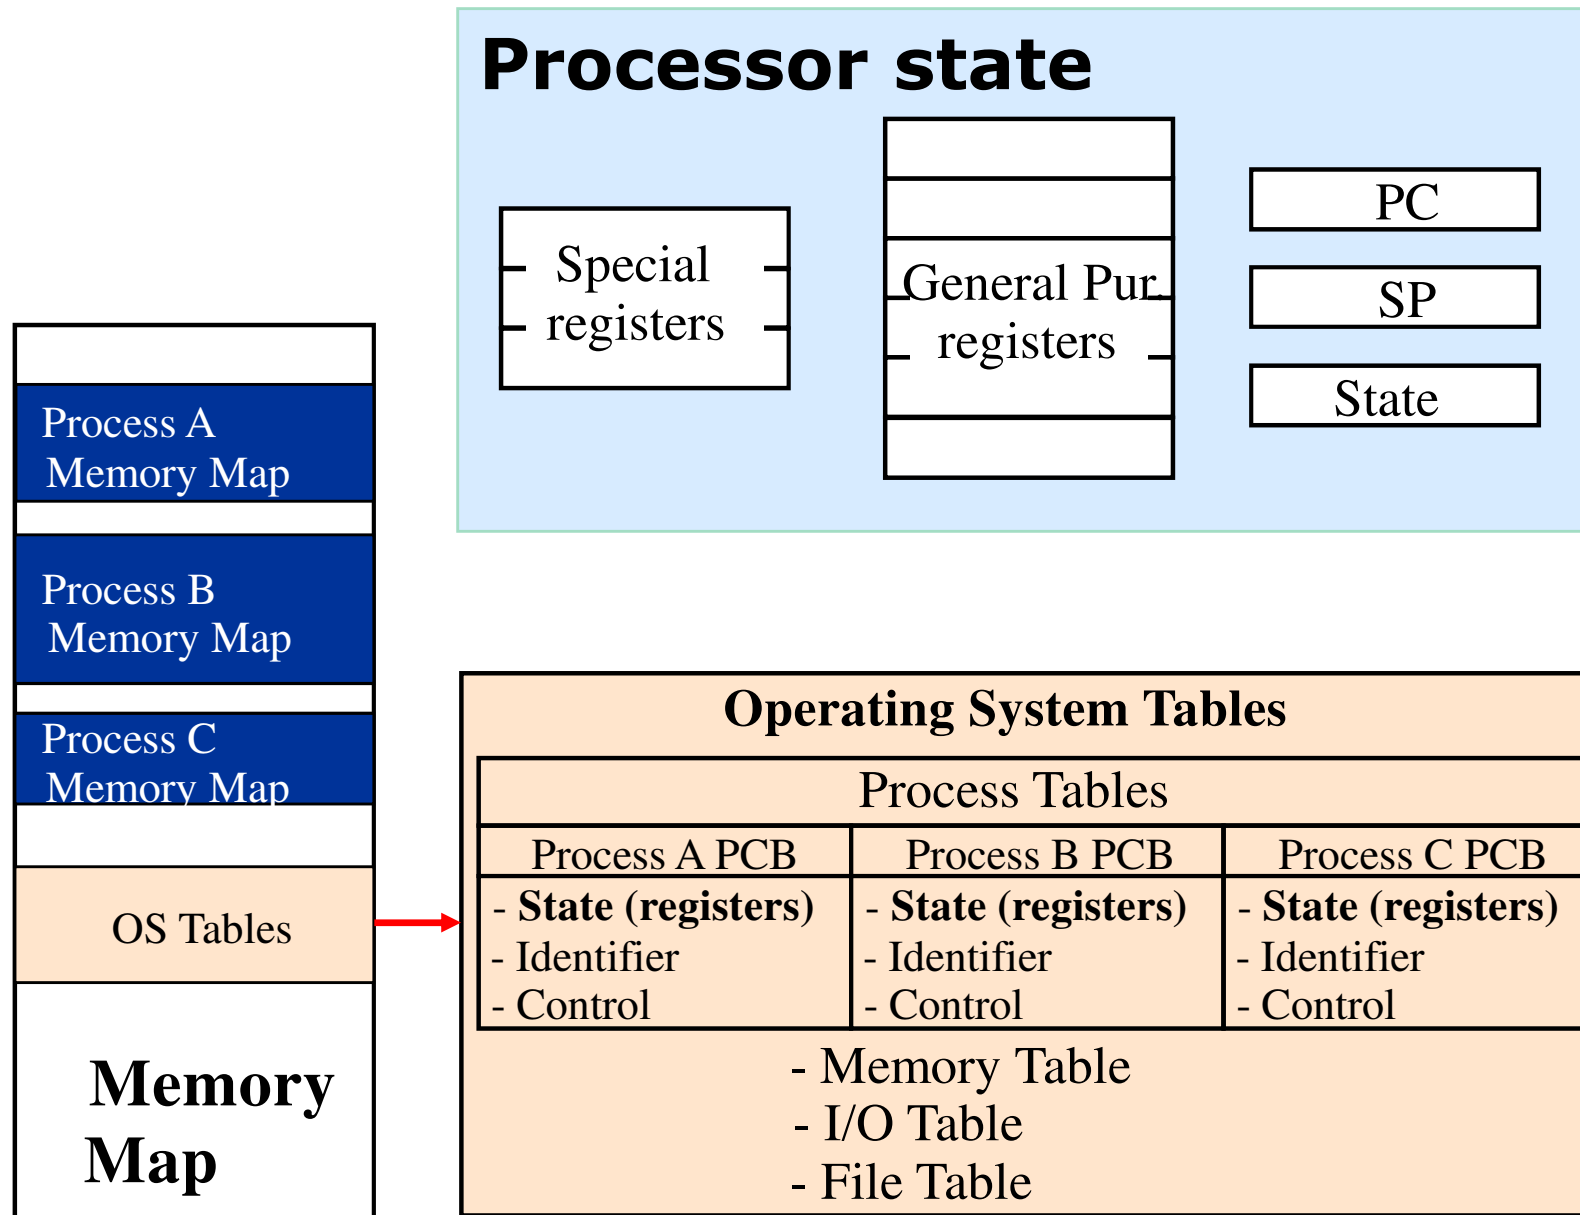


Figure 3.1 Simplified Process Control Block






Process Information




Source: Sistemas Operativos. Una visión aplicada.






What does define a process?

- State of the *processor* 
 - Data of the registers of the processor
- Core image
 - Data of the memory segments (code, data and stack)
- Information/status of each process
 - Process Control Block (PCB)

Process Image







Process organization

■ Processes organization

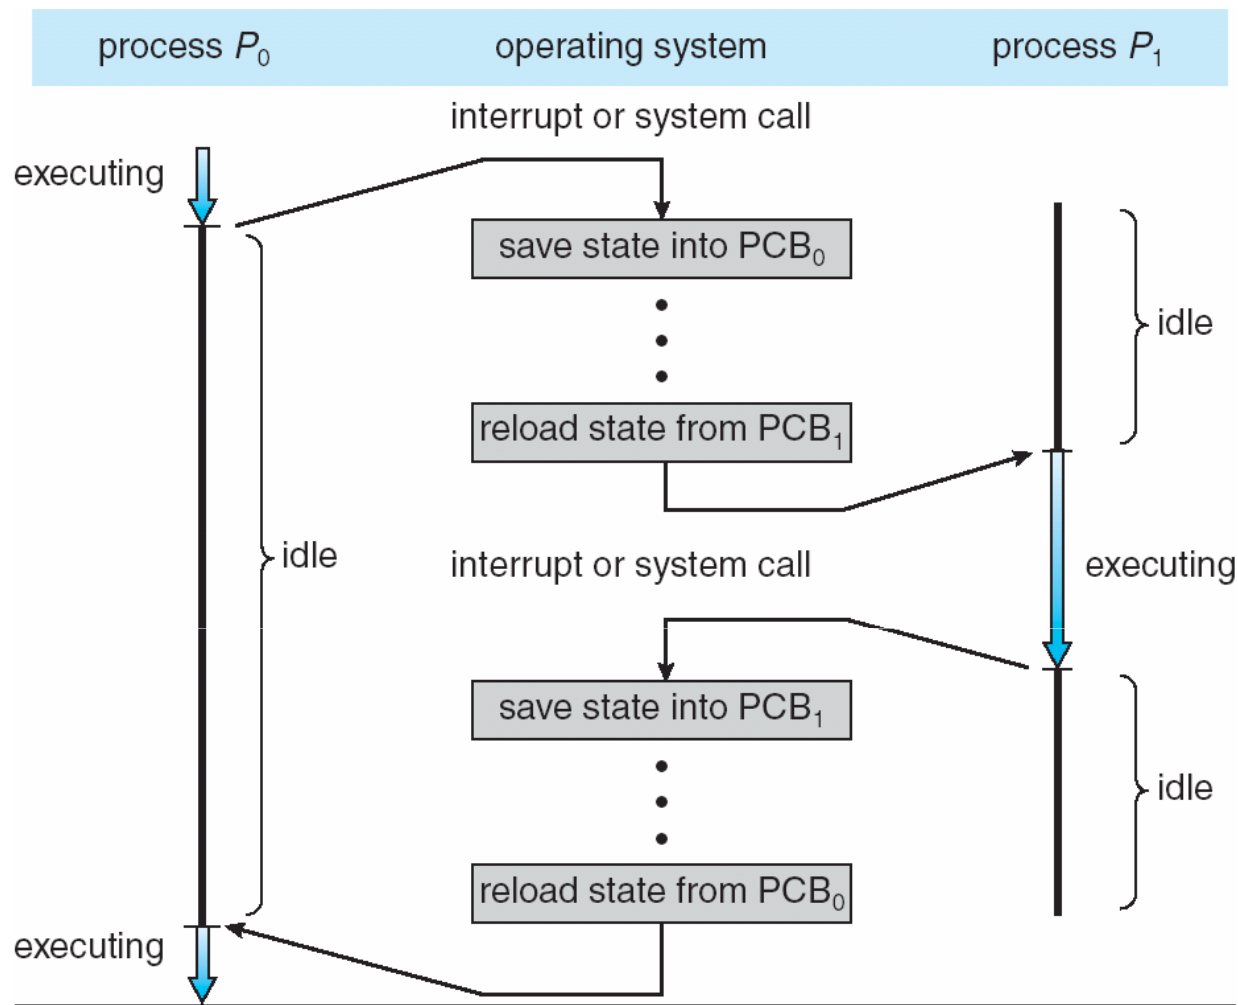
- Process lists of a certain type
 - ▶ Run: as many processes as processors
 - ▶ Ready: sorted list by the scheduler
 - ▶ Block: several non sorted lists
 - It speeds up the search of the process to wake up





CPU Switch From Process to Process


- The PCB is saved when a process is removed from the CPU and another process takes its place (context switch).





Context Switch


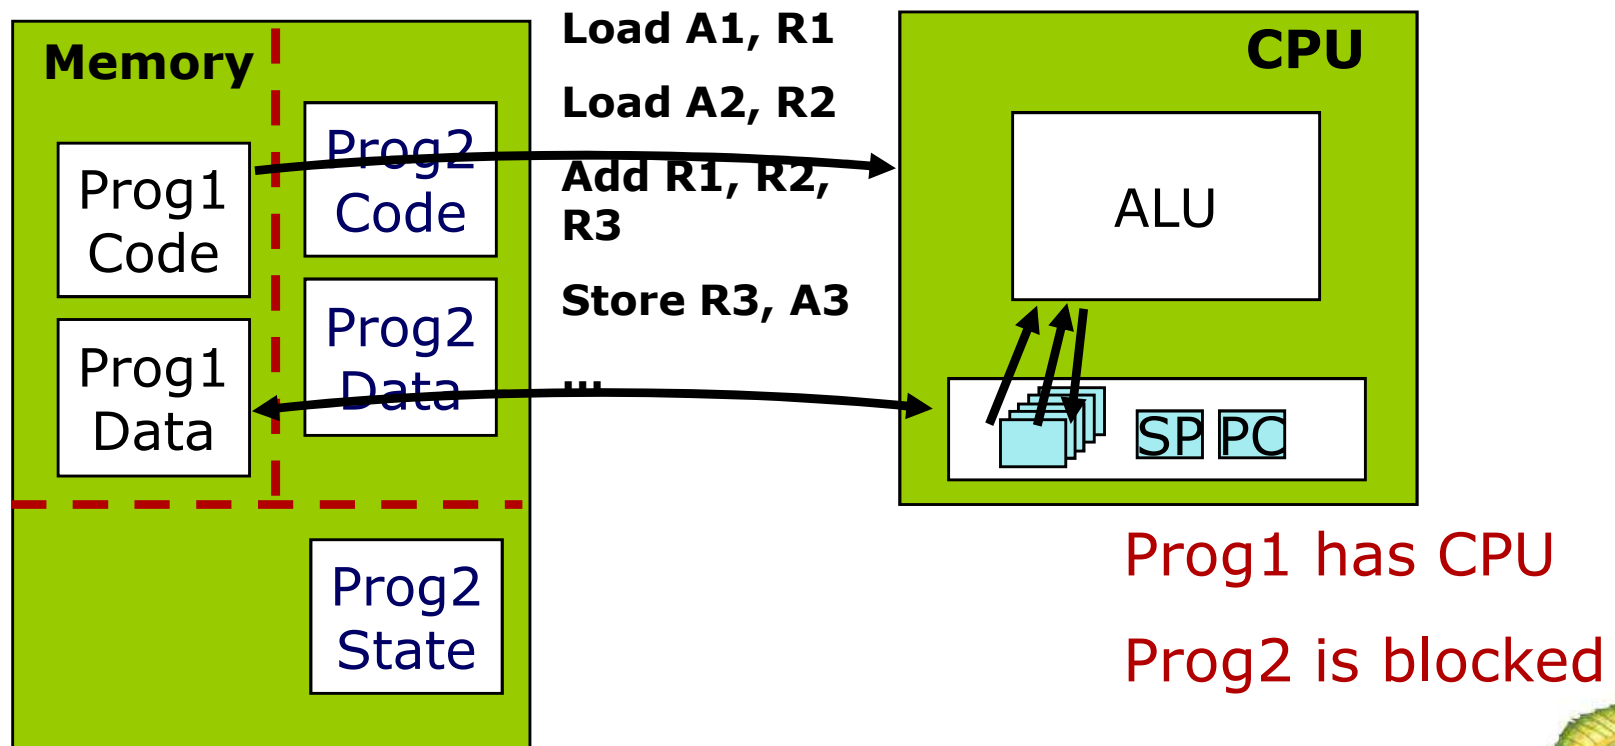
- When CPU switches to another process, the system **must save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch time is overhead; the system does no useful work while switching
 - The more complex the OS and the PCB → the longer the context switch
- Time dependent on hardware support.
 - Varies from 1 to 1000 microseconds





Switching among multiple processes



- Program instructions operate on operands in memory and (temporarily) in registers





Switching among multiple processes



- Saving all the information about a process allows a process to be *temporarily suspended* and later *resumed* from the same point





Switching among multiple processes


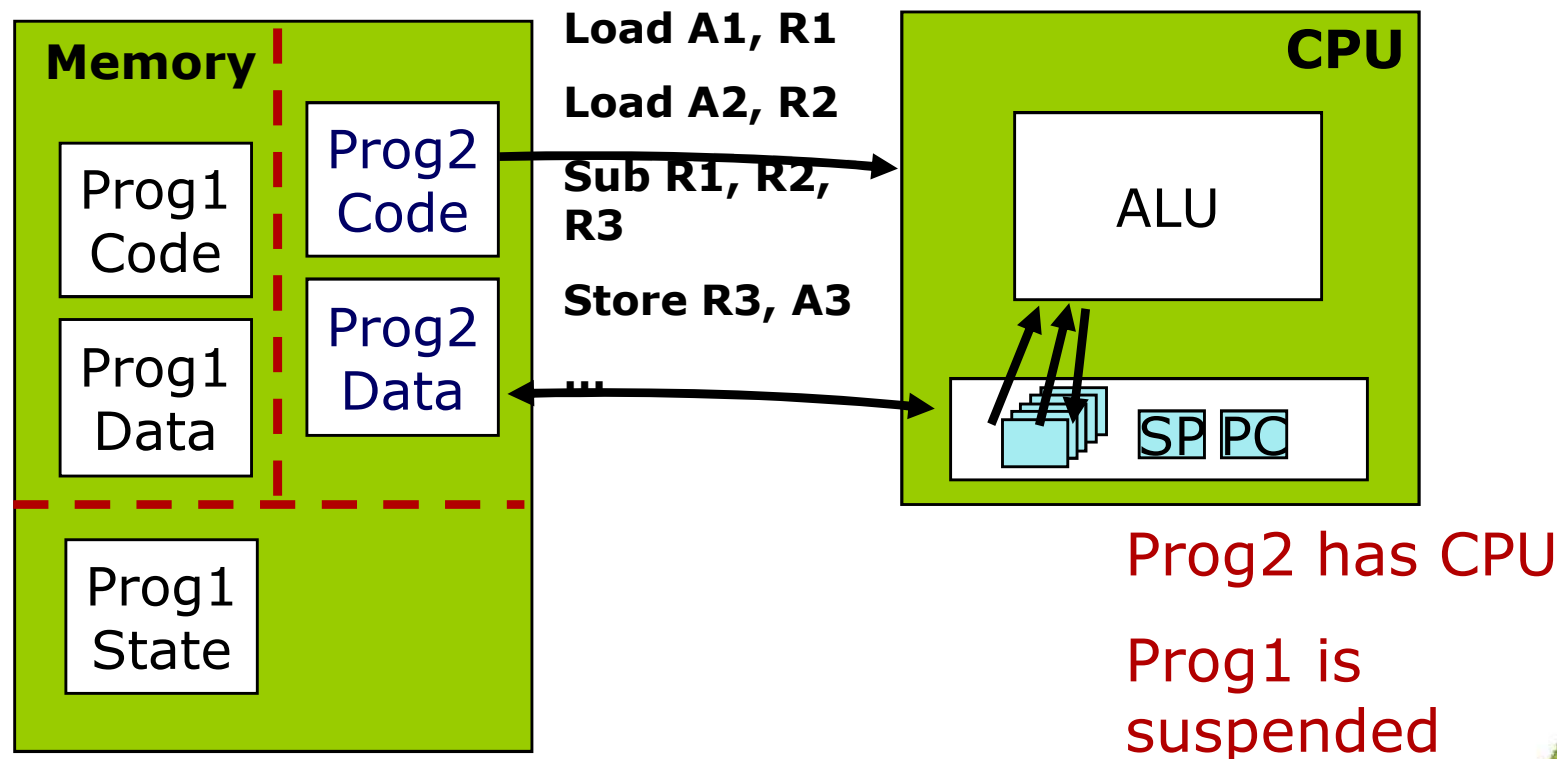
- Saving all the information about a process allows a process to be *temporarily suspended* and later *resumed*





Switching among multiple processes

- Program instructions operate on operands in memory and in registers





Contents

■ Process


- Process concept
- Process states
- Process Control Block (PCB)
- Operations on Processes

■ Threads

- Thread concept
- Libraries to create threads

■ Operating System Examples

■ CPU scheduling






How do processes get created?

Principal events that cause process creation


- System initialization
- Initiation of a batch job
- User request to create a new process
- Execution of a process creation system call from another process





Process Creation


- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
 - special system calls for communicating with and waiting for child processes
 - each process is assigned a unique identifying number or **process ID (PID)**
- Resource sharing
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution
 - Parent and children execute concurrently
 - Parent waits until children terminate





Process Creation (Cont)


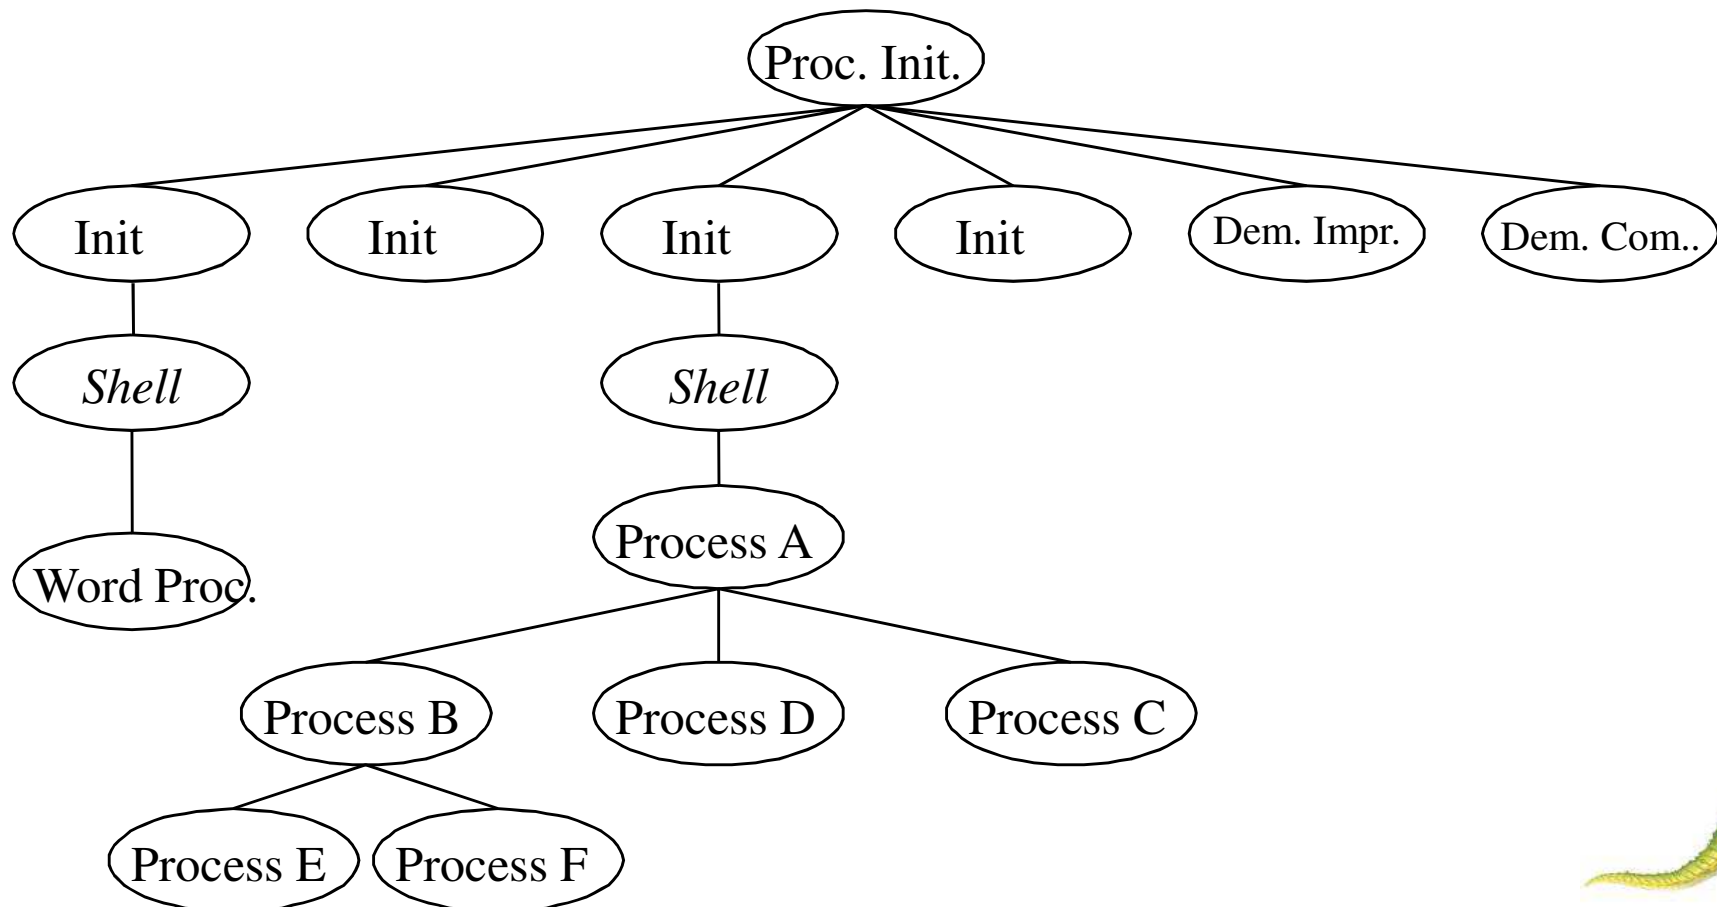
- Address space
 - Child duplicate of parent (UNIX)
 - ▶ Child has copy of parent's address space
 - Enables easy communication between the two
 - The child process' memory space is replaced with a new program which is then executed. Parent can wait for child to complete or create more processes
 - Child has a program loaded into it directly (Windows or DEC VMS)
- Child processes can create their own child processes
 - Forms a hierarchy
 - UNIX calls this a "process group"
 - Windows has no concept of process hierarchy
 - ▶ all processes are created equal





A tree of processes in Unix

- The first process is Init ()
- Init creates login daemons
- Login becomes a shell
- Using the shell user spawns new processes



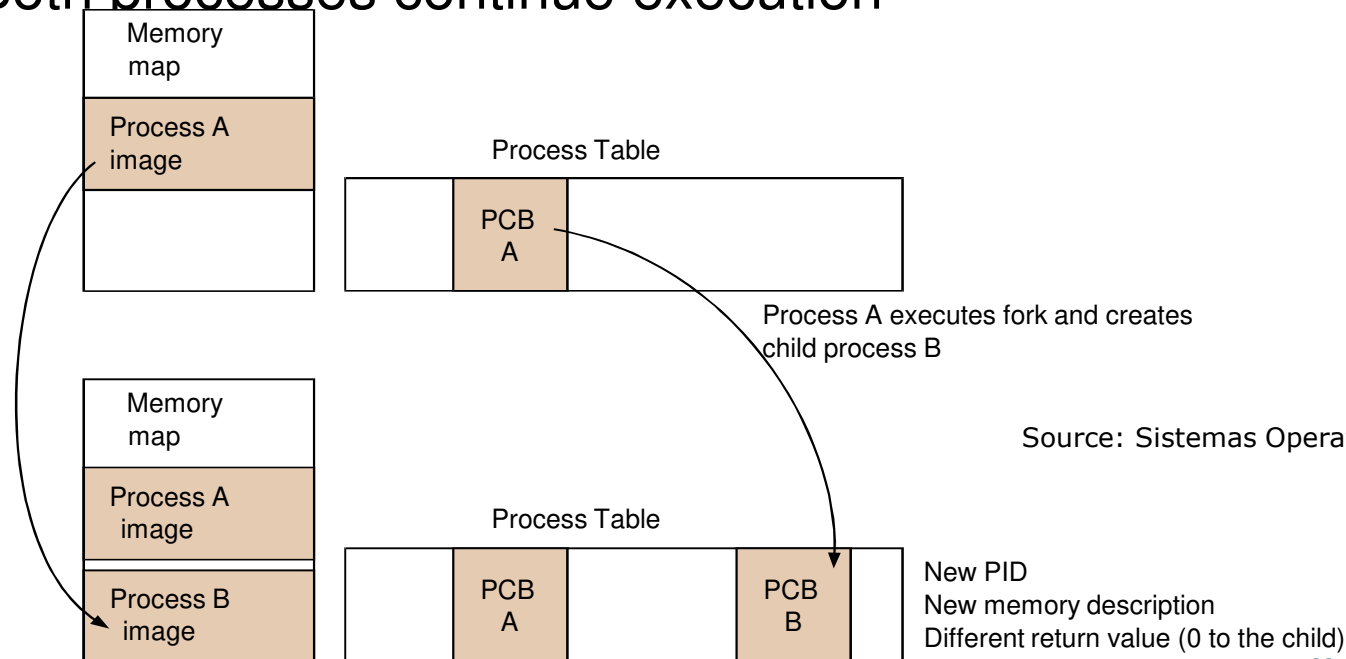


Process creation in UNIX (POSIX)


■ Process creation: `fork` and `exec`

● `pid_t fork(void)` system call

- ▶ creates new process which has a copy of the address space of the original process and returns in both processes (parent and child), but with a different return value:
 - 0 to child, pid's child to parent and -1 in error case
- ▶ Simplifies parent-child communication
- ▶ Both processes continue execution



Source: Sistemas Operativos. Una visión aplicada.

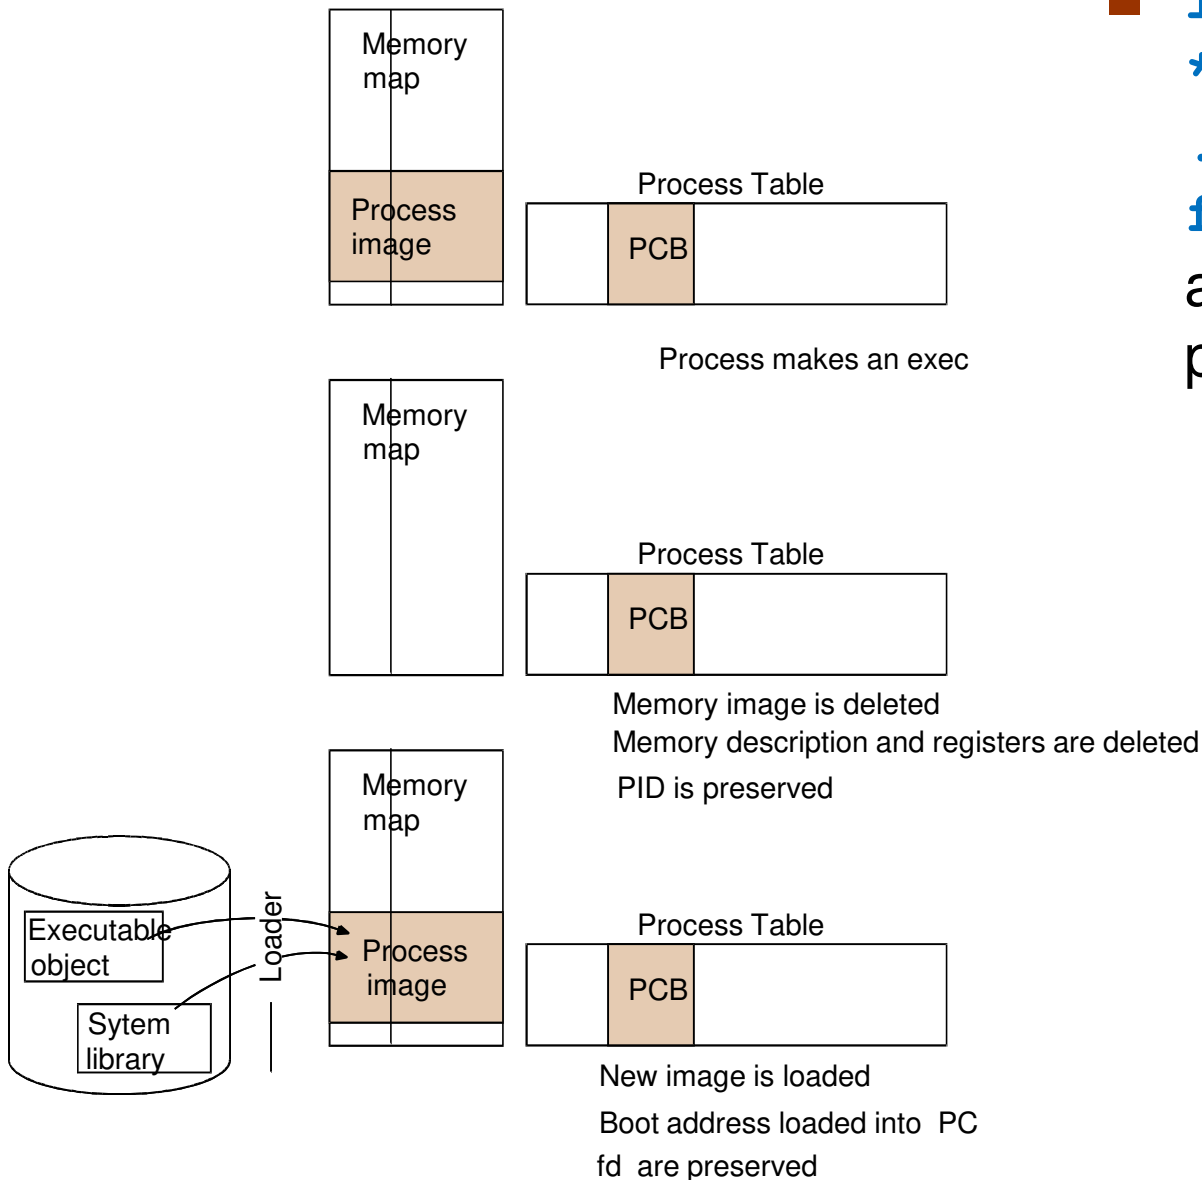





Process creation in Unix

- `int execlp(const char *file, const char *arg, ...)` system call used after a `fork()` to replace the process' address space with a new program

- Loads a binary file into memory and starts execution
- Parent can then create more children processes or issue a `wait()` system call to move itself off the ready queue until the child completes



Source: Sistemas Operativos. Una visión aplicada.





C Program Forking Separate Process


```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```






Example: process creation in UNIX

csh (pid = 22)

```
...  
  
pid = fork()  
if (pid == 0) {  
    // child..  
    ...  
    exec("ls"...);  
}  
else {  
    // parent  
    wait();  
}  
...
```






Process creation in UNIX example

csh (pid = 22)

```
...  
  
pid = fork()  
if (pid == 0) {  
    // child..  
    ...  
    exec("ls"...);  
}  
else {  
    // parent  
    wait();  
}  
...
```

csh (pid = 24)

```
...  
  
pid = fork()  
if (pid == 0) {  
    // child..  
    ...  
    exec("ls"...);  
}  
else {  
    // parent  
    wait();  
}  
...
```






Process creation in UNIX example

csh (pid = 22)

```
...  
  
pid = fork()  
if (pid == 0) {  
    // child..  
    ...  
    exec("ls"...);  
}  
else {  
    // parent  
    wait();  
}  
...
```

csh (pid = 24)

```
...  
  
pid = fork()  
if (pid == 0) {  
    // child..  
    ...  
    exec("ls"...);  
}  
else {  
    // parent  
    wait();  
}  
...
```






Process creation in UNIX example

csh (pid = 22)

```
...  
  
pid = fork()  
if (pid == 0) {  
    // child..  
    ...  
    exec("ls"...);  
}  
else {  
    // parent  
    wait();  
}  
...
```

csh (pid = 24)

```
...  
  
pid = fork()  
if (pid == 0) {  
    // child..  
    ...  
    exec("ls"...);  
}  
else {  
    // parent  
    wait();  
}  
...
```






Process creation in UNIX example

csh (pid = 22)

```
...  
  
pid = fork()  
if (pid == 0) {  
    // child..  
    ...  
    exec("ls"...);  
}  
else {  
    // parent  
    wait();  
}  
...
```

ls (pid = 24)

```
//ls program  
main() {  
    //look up dir  
    ...  
}
```





Creating a Separate Process via Windows API


```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");


    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```






Others POSIX services


- All processes have a unique process id
 - `getpid()`, `getppid()` system calls allow processes to get their information
 - ▶ `pid_t getpid(void)` —returns the process ID of the calling process
 - ▶ `pid_t getppid(void)` —returns the parent process ID of the calling process





Process Termination


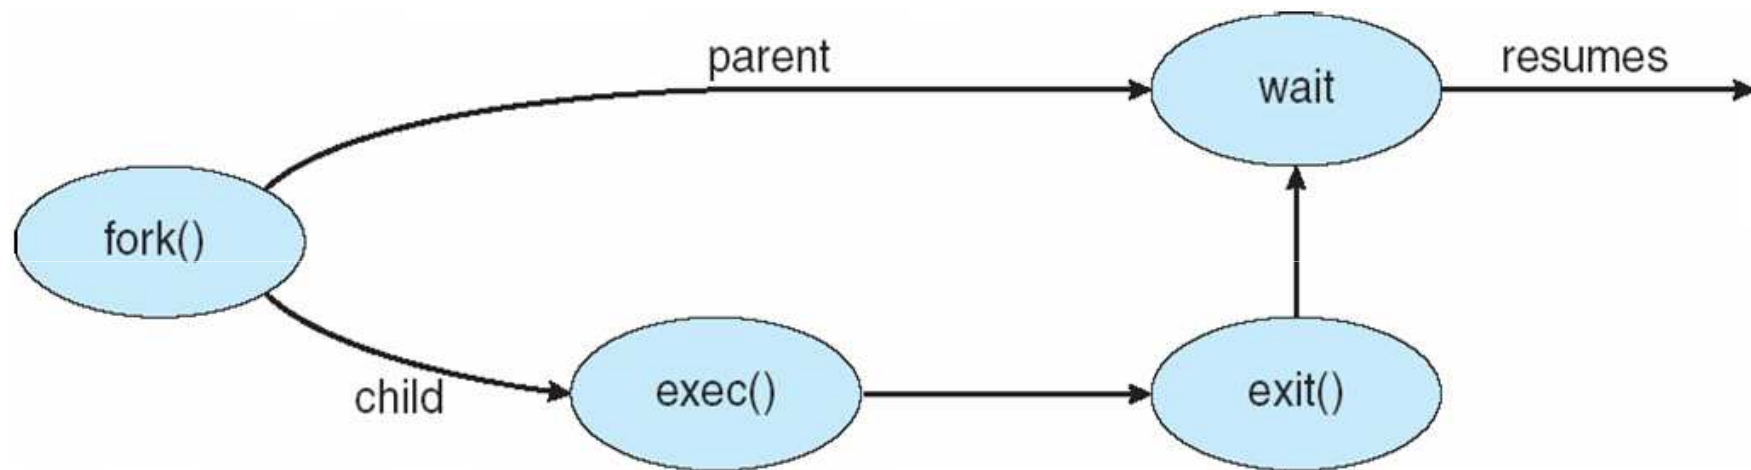
- Process executes last statement and asks the operating system to delete it by using the **exit()** system call
 - Process may return a status value to parent via **wait()**
 - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes **abort()** in UNIX
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - Parent is exiting
 - ▶ Some operating systems do not allow child to continue if its parent terminates - **cascading termination**





Process Termination in Unix

- In UNIX, a process can be terminated via the **exit** system call.
 - Parent can wait for termination of child by the **wait** system call
 - **wait** returns the process identifier of a terminated child so that the parent can tell which child has terminated






Wait() and Exit() (POSIX services)

■ `int exit (int status);`

- Arguments: code of returning to the parent
- Description:
 - ▶ It finishes process execution
 - ▶ It closes all fd
 - ▶ It releases all process resources

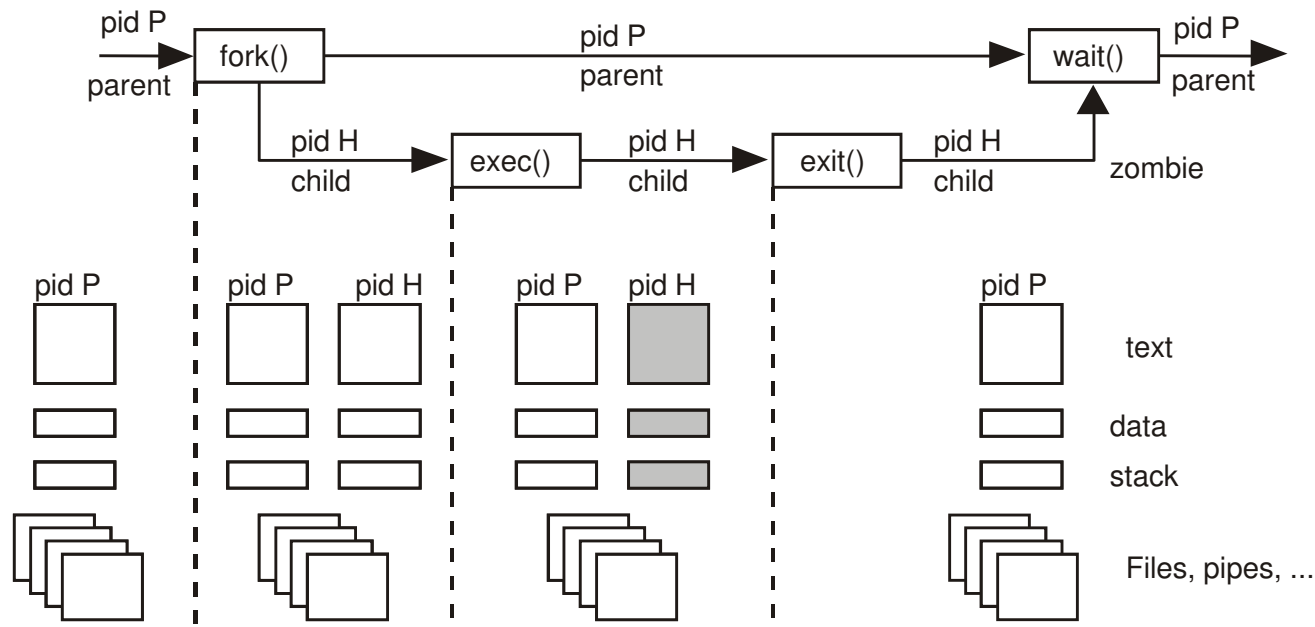





Wait() and Exit() (POSIX services)

■ `pid_t wait (int *status);`

- Arguments: returns the exit status of the child
- Description:
 - ▶ This system call suspends execution of the calling process until one of its children terminates (a parent can wait till its child finishes)
 - ▶ It returns pid and exit status of the child (-1 in error case)



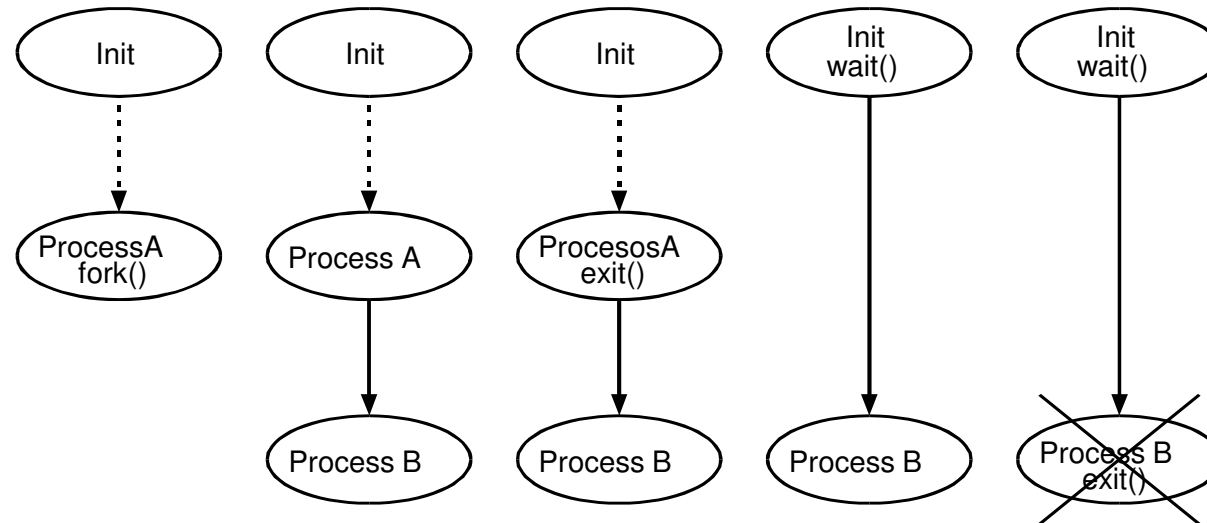
Source: Sistemas Operativos. Una visión aplicada.




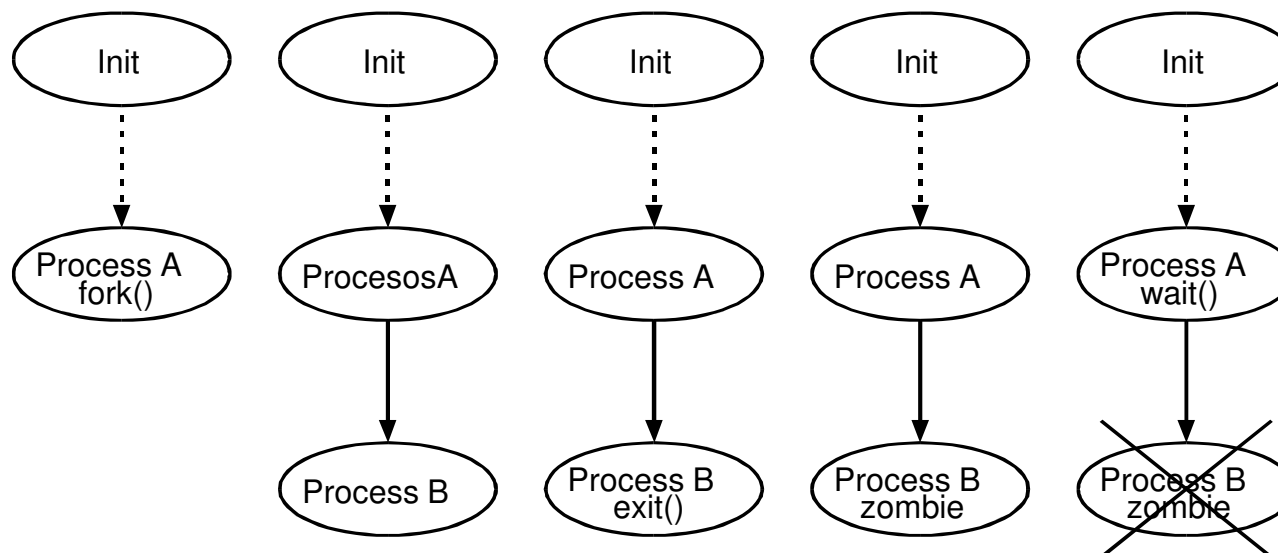



Process Termination in Unix

- If a parent terminates, all children are assigned the **init** process as their new parent




- Zombie: childrens terminate (exit) and father doesn't call to wait





Source: <https://ninefold.com/blog/2014/11/25/threads/>






Example: process creation in UNIX

■ A process launches command “ls -l”

```
#include <sys/types.h>
#include <stdio.h>

/* program executes command ls -l */
main() {
    pid_t pid;
    int status;


    pid = fork();
    if (pid == 0)    { /* child process */
        execlp("ls", "ls", "-l", NULL);
        exit(-1);
    }
    else            /* parent process */
        while (pid != wait(&status));
    exit(0);
}
```







Signal Handling

- Signals are used in UNIX systems to notify a process that a particular event has occurred
- Process termination, signaling
 - *signal()*, *kill()* system calls allow a process to be terminated or have specific signals sent to it
- The signals are process interrupts
 - Signals are sent:
 - ▶ From process to process with kill
 - ▶ From OS to a process




Source: Sistemas Operativos. Una visión aplicada.





Signal Handling


- Types of signals (kill -l shows the list)
 - **SIGKILL** -9- (Kill : terminate immediately)
 - **SIGTERM** -15- (Termination : request to terminate)
 - **SIGHLD** -17- (Child process terminated, stopped or continued)
 - **SIGCONT** -18- (Continue if stopped)
 - **SIGSTOP** -19- (Stop executing temporarily)
 - **SIGTTIN** -21. (Background process attempting to read from tty)
 - **SIGTTOU** -22- (Background process attempting to write to tty)





Signal Handling


- **Signals** are used in UNIX systems to notify a process that a particular event has occurred.
- A **signal handler** is used to process signals
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Signal is handled by one of two signal handlers:
 1. default
 2. user-defined
- Every signal has **default handler** that kernel runs when handling signal
 - **User-defined signal handler** can override default





Signal Handling in Unix


- Signal handlers can be installed with the `signal()` system call.
 - When the signal is intercepted, the signal handler is invoked
- If a signal handler is not installed for a particular signal, the default handler is used.
- The process can also specify two default behaviors, without creating a handler:
 - ignore the signal (SIG_IGN)
 - use the default signal handler (SIG_DFL)
- The `sigprocmask()` call can be used to block and unblock delivery of signals.





Handling Signals (POSIX)

- `int kill(pid_t pid, int sig)`
 - It sends to the process “pid” the signal “sig”
- `int sigaction(int sig, struct sigaction *act, struct sigaction *oact)`
 - It allows to specify the action to be taken when the signal “sig” is received
- `int pause(void)`
 - It blocks the process until the reception of a signal
- `unsigned int alarm(unsigned int seconds)`
 - It generates the reception of signal SIGALARM after “seconds” seconds
- `sigprocmask(int how, const sigset_t *set, sigset_t *oact)`
 - It is used to explore or modify the signal mask of a process





Win32 services

- To create a process
 - `BOOL CreateProcess (...);`
- To finish the execution of a process
 - `VOID ExitProcess (UINT nExitCode);`
- To obtain the end code of a process
 - `BOOL GetExitCodeProcess (HANDLE hProcess, LPDWORD lpdwExitCode);`
- To finish the execution of another process
 - `BOOL TerminateProcess (HANDLE hProcess, UINT uExitCode);`
- To wait till the end of a process
 - `DWORD WaitForSingleObject (HANDLE hObject, DWORD dwTimeout);`
 - `DWORD WaitForMultipleObjects (DWORD cObjects, LPHANDLE lphObjects, BOOL fWaitAll, DWORD dwTimeout);`





Summary (POSIX & Win32 APIs)

- Process identification
 - `getpid`, `GetCurrentProcessId`
- Process environment variables
 - `getenv`, `GetEnvironmentStrings`
 - Path, home directory, working directory, personal directory
- Process creation: `fork`, `CreateProcess`
 - Program transformation: `exec`
- Waiting for termination of a process: `wait`, `WaitForSingleObject`, `WaitForMultipleObjects`
- Process termination: `exit`, `ExitProcess`
- Send signal to process: `kill`, `TerminateProcess`

